

Andrew O. Finley and Jeffrey W. Doser

Forestry 472: Ecological Monitoring and Data Analysis



Contents

List of Tables	3
List of Figures	5
Preface	7
1 Data	9
1.1 FEF Tree Biomass Data Set	9
1.2 FACE Experiment Data Set	9
1.3 PEF Inventory and LiDAR Data Set	11
1.4 Zurichberg Forest inventory data set	13
1.5 Looking Forward	13
1.6 How to Learn (The Most Important Section in This Book!)	14
2 Introduction to R and RStudio	17
2.1 Obtaining and Installing R	17
2.2 Obtaining and Installing RStudio	18
2.3 Using R and RStudio	18
2.3.1 R as a calculator	19
2.3.2 Basic descriptive statistics and graphics in R	21
2.3.3 Simple linear regression in R	23
2.4 How to Learn	24
2.5 Getting help	25
2.6 Workspace, working directory, and keeping organized	26
2.7 Quality of R code	27
2.7.1 Naming Files	27
2.7.2 Naming Variables	27
2.7.3 Syntax	28
3 Scripts, R Markdown, and Reproducible Research	29
3.1 Scripts in R	30
3.2 R Markdown	34
3.2.1 Creating and processing R Markdown documents	37
3.2.2 Text: Lists and Headers	37
3.2.3 Code Chunks	38
3.2.4 Output formats other than HTML	42
3.2.5 LaTeX, knitr, and bookdown	42

4	Data Structures	43
4.1	Vectors	44
4.1.1	Types, Conversion, and Coercion	45
4.1.2	Accessing Specific Elements of Vectors	48
4.2	Factors	50
4.3	Names of objects in R	52
4.4	Missing Data, Infinity, etc.	53
4.4.1	Infinity and NaN	54
4.5	Data Frames	55
4.5.1	Accessing specific elements of data frames	57
4.6	Lists	59
4.6.1	Accessing specific elements of lists	61
4.7	Subsetting with Logical Vectors	63

List of Tables

1.1	A subset of the tree biomass data from the FEF.	10
1.2	A small portion of the FACE experiment data set	11
1.3	A small portion of the PEF inventory data set	12
4.1	Dimension and Type Content of Base Data Structures in R. .	43



List of Figures

1.1	Surface of LiDAR energy returns at 12 m above the ground, forest inventory plot locations, and management unit boundaries on the PEF.	13
1.2	Location and species of all trees in the Zurichberg Forest. . .	14
2.1	The Rstudio IDE.	19
2.2	xkcd: Code Quality	27
3.1	A script window in RStudio	34
3.2	Example R Markdown Input and Output	36
3.3	Producing Lists in R Markdown	38
3.4	Headers and Some LaTeX in R Markdown	39
3.5	Other useful LaTeX symbols and expressions in R Markdown	40
3.6	Output of Example R Markdown	41



Preface

This text is an introduction to data sciences for Forestry and Environmental students. Understanding and responding to current environmental challenges requires strong quantitative and analytical skills. There is a pressing need for professionals with data science expertise in this data rich era. The McKinsey Global Institute¹ predicts that “by 2018, the United States alone could face a shortage of 140,000 to 190,000 people with deep analytical skills as well as 1.5 million managers and analysts with the know-how to use the analysis of big data to make effective decisions”. The Harvard Business Review dubbed *data scientist* “The Sexiest Job of the 21st Century”². This need is not at all confined to the tech sector, as forestry professionals are increasingly asked to assume the role of *data scientists* and *data analysts* given the rapid accumulation and availability of environmental data (see, e.g. Schimel and Keller (2015)). Thomson Nguyen’s talk³ on the difference between a data scientist and a data analyst is very interesting and contains elements relevant to the aim of this text. This aim is to give you the opportunity to acquire the tools needed to become an environmental data analyst. Following Bravo et al. (2016) a *data analyst* has the ability to make appropriate calculations, convert data to graphical representation, interpret the information presented in graphical or mathematical forms, and make judgements or draw conclusions based on the quantitative analysis of data.

¹http://www.mckinsey.com/insights/business_technology/big_data_the_next_frontier_for_innovation

²<https://hbr.org/2012/10/data-scientist-the-sexiest-job-of-the-21st-century>

³www.import.io/post/data-scientists-vs-data-analysts-why-the-distinction-matters



1

Data

1.1 FEF Tree Biomass Data Set

When thinking about data, we might initially have in mind a modest-sized and uncomplicated data set that serves a fairly specific purpose. For example, in forestry it is convenient to have a mathematical formula that relates a tree's diameter (or some other easily measured attribute) to stem or total biomass (i.e. we cannot directly measure tree biomass without destructive sampling). When coupled with forest inventory data, such formulas provide a means to estimate forest biomass across management units or entire forest landscapes. A data set used to create such formulas includes felled tree biomass by tree component for four hardwood species of the central Appalachians sampled on the Fernow Experimental Forest¹ (FEF), West Virginia Wood et al. (2016). A total of 88 trees were sampled from plots within two different watersheds on the FEF. Hardwood species sampled include *Acer rubrum*, *Betula lenta*, *Liriodendron tulipifera*, and *Prunus serotina*, all of which were measured in the summer of 1991 and 1992. Data include tree height, diameter, as well as green and dry weight of tree stem, top, small branches, large branches, and leaves. Table 1.1 shows a subset of these data

The size of this dataset is relatively small, there are no missing observations, the variables are easily understood, etc.

1.2 FACE Experiment Data Set

We often encounter data gleaned from highly structured and complex experiments. Such data typically present challenges in organization/storage, exploratory data analysis (EDA), statistical analysis, and interpretation of analysis results. An example data set comes from the Aspen Free-Air Carbon Diox-

¹<http://www.nrs.fs.fed.us/ef/locations/wv/fernow>

TABLE 1.1: A subset of the tree biomass data from the FEF.

species	dbh_in	height_ft	stem_green_kg	leaves_green_kg
Acer rubrum	6.0	48.0	92.2	16.1
Acer rubrum	6.9	48.0	102.3	12.9
Acer rubrum	6.4	48.0	124.4	16.5
Acer rubrum	6.5	49.0	91.7	12.0
Acer rubrum	7.2	51.0	186.2	22.4
Acer rubrum	3.1	40.0	20.8	0.9
Acer rubrum	2.0	30.5	5.6	1.0
Acer rubrum	4.1	50.0	54.1	6.1
Acer rubrum	2.4	28.0	10.2	2.5
Acer rubrum	2.7	40.4	20.2	1.6

ide Enrichment² (FACE) Experiment conducted from 1997-2009 on the Harshaw Experimental Forest³ near Rhinelander, Wisconsin. The Aspen FACE Experiment was a multidisciplinary study that assessed the effects of increasing tropospheric ozone and carbon dioxide concentrations on the structure and functioning of northern forest ecosystems. The design provided the ability to assess the effects of these gasses alone (and in combination) on many ecosystem attributes, including growth, leaf development, root characteristics, and soil carbon. The data set considered here comprises annual tree height and diameter measurements from 1997 to 2008 for *Populus tremuloides*, *Acer saccharum*, and *Betula papyrifera* grown within twelve 30 meter diameter rings in which the concentrations of tropospheric ozone and carbon dioxide were controlled Kubiske (2013). Because there was no confinement, there was no significant change in the natural, ambient environment other than elevating these trace gas concentrations. Although the basic individual tree measurements are similar to those in the FEF data set we saw in Section 1.1, (i.e., height and diameter), the study design specifies various tree species clones, varying gas treatments, and treatment replicates. Further, because these are longitudinal data, (measurements were recorded over time) the data set presents many missing values as a result of tree mortality. Table 1.2 contains the first five records as well as 5 more randomly selected records in the data set. Here, a row identifies each tree's experimental assignment, genetic description, and growth over time.

Notice that several height measurements in 2008 contain missing data. If all year measurements were shown, we would see much more missing data. Also, notice that this data set is substantially larger than the FEF data set with 912 rows and 39 columns of data in the full data set.

²http://www.nrs.fs.fed.us/disturbance/climate_change/face

³<http://www.nrs.fs.fed.us/ef/locations/wi/rhineland/>

TABLE 1.2: A small portion of the FACE experiment data set

	Rep	Treat	SPP	Clone	ID..	X1997_Height	X2008_Height
1	1	1	B	B	4360	51.0	632
2	1	1	A	216	4359	58.0	742
3	1	1	B	B	4358	24.0	916
4	1	1	A	216	4357	58.0	981
5	1	1	B	B	4356	41.0	914
183	1	3	B	B	5017	40.0	NA
625	3	1	B	B	6853	55.5	936
835	3	3	B	B	7573	48.0	NA
259	1	4	B	B	5327	48.0	659
96	1	2	A	216	4697	27.0	NA

1.3 PEF Inventory and LiDAR Data Set

Coupling forest inventory with remotely sensed Light Detection and Ranging (LiDAR) data sets using regression models offers an attractive approach to mapping forest variables at stand, regional, continental, and global scales. LiDAR data have shown great potential for use in estimating spatially explicit forest variables over a range of geographic scales (Asner et al., 2009), (Babcock et al., 2013), (Finley et al., 2011), (Næsset, 2011), (Neigh et al., 2013). Encouraging results from these and many other studies have spurred massive investment in new LiDAR sensors, sensor platforms, as well as extensive campaigns to collect field-based calibration data.

Much of the interest in LiDAR based forest variable mapping is to support carbon monitoring, reporting, and verification (MRV) systems, such as defined by the United Nations Programme on Reducing Emissions from Deforestation and Forest Degradation⁴ (UN-REDD) and NASA's Carbon Monitoring System⁵ (CMS) (Le Toan et al., 2011), (Ometto et al., 2014). In these, and similar initiatives, AGB is the forest variable of interest because it provides a nearly direct measure of forest carbon (i.e., carbon comprises $\sim 50\%$ of wood biomass, West (2004)). Most efforts to quantify and/or manage forest ecosystem services (e.g., carbon, biodiversity, water) seek high spatial resolution wall-to-wall data products such as gridded maps with associated measures of uncertainty, e.g., point and associated credible intervals (CIs) at the pixel level. In fact several high profile international initiatives include language con-

⁴<http://www.un-redd.org>

⁵<http://carbon.nasa.gov>

TABLE 1.3: A small portion of the PEF inventory data set

	MU	plot	easting	northing	biomass.mg.ha	stocking.stems.ha
118	17	24	530304	4965983	112.96	2325
403	31	11	530575	4964959	NA	NA
539	8	23	530004	4967094	71.05	10927
167	19	63	530436	4965217	NA	NA
62	14	21	530218	4966445	NA	NA
410	31	32	530657	4964999	NA	NA
308	27	31	530449	4965815	134.35	7872
471	6	13	529560	4967220	30.33	3016
556	9	14	529601	4966363	140.00	1284
65	14	24	530339	4966652	163.52	433

cerning the level of spatially explicit acceptable error in total forest carbon estimates, see, e.g., [UN-REDD \(2009\)](#) and [UNFCCC \(2015\)](#).

Here, we consider a data set collected on the Penobscot Experimental Forest⁶ (PEF) in Bradley and Eddington, Maine. The dataset comprises LiDAR waveforms collected with the Laser Vegetation Imaging Sensor⁷ (LVIS) and several forest variables measured on a set of 589 georeferenced forest inventory plots. The LVIS data were acquired during the summer of 2003. The LVIS instrument, an airborne scanning LiDAR with a 1064 nm laser, provided 12,414 LiDAR pseudo-waveform signals within the PEF. For each waveform, elevations were converted to height above the ground surface and interpolated at 0.3 m intervals. Figure 1.1 shows PEF LiDAR energy returns at 12 m above the ground, forest inventory plot locations, and management unit boundaries. The forest inventory data associated with each plot were drawn from the PEF’s database of several on-going, long-term silvicultural experiments (see [Kenefic et al. \(2015\)](#)). Below we provide a plot containing the geographic coordinates, biomass (mg/ha), basal area (m²/ha), stocking (trees/ha), diameter class (cm), and management unit. Table 1.3 shows a subset of data for 10 randomly selected plots (where each row records plot measurements) in the forest inventory data set.

⁶www.fs.fed.us/ne/durham/4155/penobsco.htm

⁷<http://lvis.gsfc.nasa.gov>

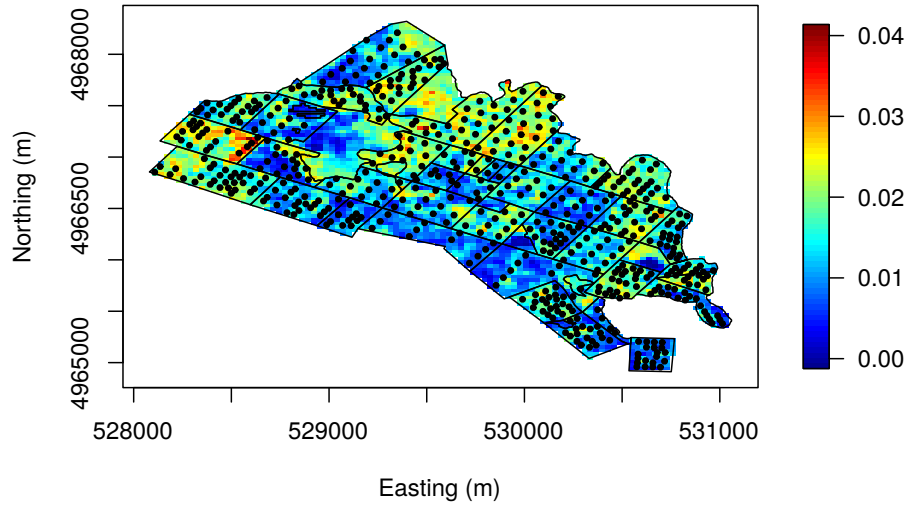


FIGURE 1.1: Surface of LiDAR energy returns at 12 m above the ground, forest inventory plot locations, and management unit boundaries on the PEF.

1.4 Zurichberg Forest inventory data set

Measuring tree diameter and height is a time consuming process. This fact makes the Zurichberg Forest inventory data set a rare and impressive investment. These data comprise a complete enumeration of the 589 trees in the Zurichberg Forest, including species, diameter at breast height, basal area, and volume. The stem map colored by species is shown in Figure 1.2.

1.5 Looking Forward

The four examples above illustrate a variety of data sets that might be encountered in practice, and each provides its own challenges. For the FACE data, the challenges are more statistical in nature. Complications could arise related to the complex study design and how that design might affect methods of analysis and conclusions drawn from the study. The other data sets present different challenges, such as how to:

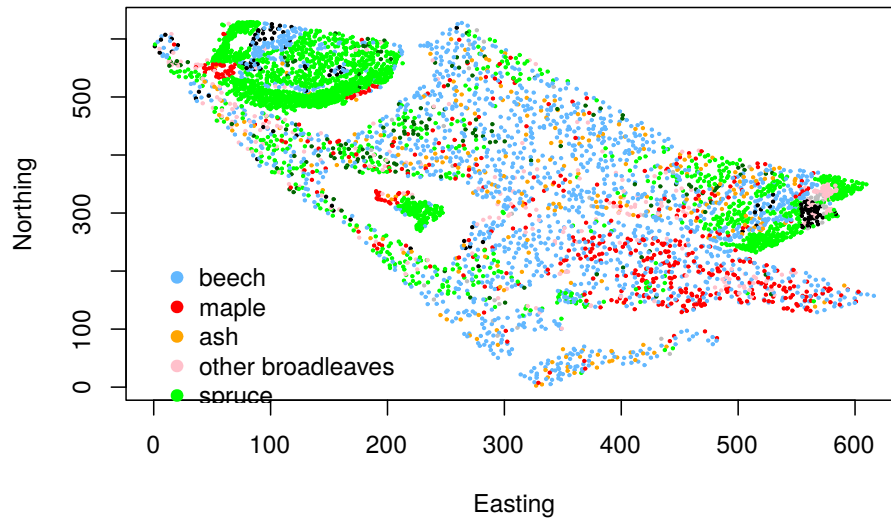


FIGURE 1.2: Location and species of all trees in the Zurichberg Forest.

1. Develop biomass equations suitable for population inference from the FEF's small sample of 88 trees
2. Work with spatially indexed data in the case of the PEF and Zurichberg inventory data
3. Effectively and efficiently process the PEF's high-dimensional LiDAR signal data for use in predictive models of forest variables.

This book and associated material introduce tools to tackle some of the challenges in working with real data sets within the context of the R statistical system. We will focus on important topics such as

- Obtaining and manipulating data
- Summarizing and visualizing data
- Communicating findings about data that support reproducible research
- Programming and writing functions
- Working with specialized data structures, e.g., spatial data and databases

1.6 How to Learn (The Most Important Section in This Book!)

There are several ways to engage with the content of this book and associated materials.

One way is not to engage at all. Leave the book closed on a shelf and do something else with your time. That may or may not be a good life strategy, depending on what else you do with your time, but you won't learn much from the book!

Another way to engage is to read through the book "passively", reading all that's written but not reading the book with R open on your computer, where you could enter the R commands from the book. With this strategy you'll probably learn more than if you leave the book closed on a shelf, but there are better options.

A third way to engage is to read the book while you're at a computer with R, and to enter the R commands from the book as you read about them. You'll likely learn more this way.

A fourth strategy is even better. In addition to reading and entering the commands given in the book, you think about what you're doing, and ask yourself questions (which you then go on to answer). For example, after working through some R code computing the logarithm of positive numbers you might ask yourself, "What would R do if I asked it to calculate the logarithm of a negative number? What would R do if I asked it to calculate the logarithm of a really large number such as one trillion?" You could explore these questions easily by just trying things out in the R Console window.

If your goal is to maximize the time you have to binge-watch *Stranger Things* Season 2 on Netflix, the first strategy may be optimal. But if your goal is to learn a lot about computational tools for data science, the fourth strategy is probably going to be best.



2

Introduction to R and RStudio

Various statistical and programming software environments are used in data science, including R, Python, SAS, C++, SPSS, and many others. Each has strengths and weaknesses, and often two or more are used in a single project. This book focuses on R for several reasons:

1. R is free
2. It is one of, if not the, most widely used software environments in data science
3. R is under constant and open development by a diverse and expert core group
4. It has an incredible variety of contributed packages
5. A new user can (relatively) quickly gain enough skills to obtain, manage, and analyze data in R

Several enhanced interfaces for R have been developed. Generally such interfaces are referred to as *integrated development environments (IDE)*. These interfaces are used to facilitate software development. At minimum, an IDE typically consists of a source code editor and build automation tools. We will use the RStudio IDE, which according to its developers “is a powerful productive user interface for R.”¹ RStudio is widely used, it is used increasingly in the R community, and it makes learning to use R a bit simpler. Although we will use RStudio, most of what is presented in this book can be accomplished in R (without an added interface) with few or no changes.

2.1 Obtaining and Installing R

It is simple to install R on computers running Microsoft Windows, macOS, or Linux. For other operating systems users can compile the source code directly.²

¹<http://www.rstudio.com/>

²Windows, macOS, and Linux users also can compile the source code directly, but for most it is a better idea to install R from already compiled binary distributions.

Here is a step-by-step guide to installing R for Microsoft Windows.³ macOS and Linux users would follow similar steps.

1. Go to <http://www.r-project.org/>
2. Click on the CRAN link on the left side of the page
3. Choose one of the mirrors.⁴
4. Click on Download R for Windows
5. Click on base
6. Click on Download R 3.5.0 for Windows
7. Install R as you would install any other Windows program

2.2 Obtaining and Installing RStudio

You must install R prior to installing RStudio. RStudio is also simple to install:

1. Go to <http://www.rstudio.com>
2. Click on the link RStudio under the Products tab, then select the Desktop option
3. Click on the Desktop link
4. Choose the DOWNLOAD RSTUDIO DESKTOP link in the Open Source Edition column
5. On the ensuing page, click on the Installer version for your operating system, and once downloaded, install as you would any program

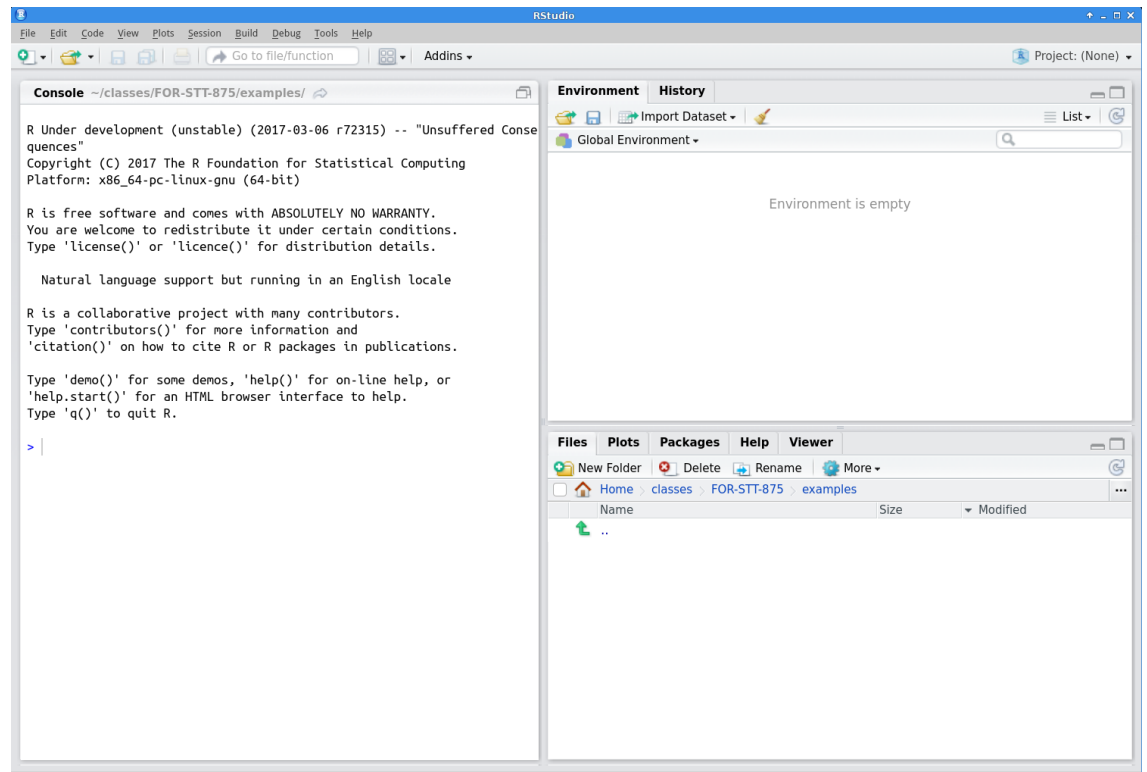
2.3 Using R and RStudio

Start RStudio as you would any other program in your operating system. For example, under Microsoft Windows use the Start Menu or double click on the shortcut on the desktop (if a shortcut was created in the installation process). A (rather small) view of RStudio is displayed in Figure 2.1.

Initially the RStudio window contains three smaller windows. For now our main focus will be the large window on the left, the Console window, in which

³New versions of R are released regularly, so the version number in Step 6 might be different from what is listed below.

⁴The <http://cran.rstudio.com/> mirror is usually fast. Otherwise choose a mirror in Michigan.

**FIGURE 2.1:** The Rstudio IDE.

R statements are typed. The next few sections give simple examples of the use of R. In these sections we will focus on small and non-complex data sets, but of course later in the book we will work with much larger and more complex sets of data. Read these sections at your computer with R running, and enter the R commands there to get comfortable using the R console window and RStudio.

2.3.1 R as a calculator

R can be used as a calculator. Note that `#` is the comment character in R, so R ignores everything following this character. Also, you will see that R prints `[1]` before the results of each command. Soon we will explain its relevance, but ignore this for now. The command prompt in R is the greater than sign `>`.

```
> 34 + 20 * sqrt(100)  ## +,-,*,/ have the expected meanings
```

```
[1] 234
```

```
> exp(2)  ##The exponential function
```

```
[1] 7.389
```

```
> log10(100)  ##Base 10 logarithm
```

```
[1] 2
```

```
> log(100)  ##Base e logarithm
```

```
[1] 4.605
```

```
> 10^log10(55)
```

```
[1] 55
```

Most functions in R can be applied to vector arguments rather than operating on a single argument at a time. A *vector* is a data structure that contains elements of the same data type (i.e. integers).

```
> 1:25  ##The integers from 1 to 25
```

```
[1]  1  2  3  4  5  6  7  8  9 10 11 12 13 14 15 16 17  
[18] 18 19 20 21 22 23 24 25
```

```
> log(1:25)  ##The base e logarithm of these integers
```

```
[1] 0.0000 0.6931 1.0986 1.3863 1.6094 1.7918 1.9459  
[8] 2.0794 2.1972 2.3026 2.3979 2.4849 2.5649 2.6391  
[15] 2.7081 2.7726 2.8332 2.8904 2.9444 2.9957 3.0445  
[22] 3.0910 3.1355 3.1781 3.2189
```

```
> 1:25 * 1:25  ##What will this produce?
```

```
[1]  1  4  9 16 25 36 49 64 81 100 121 144  
[13] 169 196 225 256 289 324 361 400 441 484 529 576  
[25] 625
```

```
> 1:25 * 1:5 ##What about this?
```

```
[1] 1 4 9 16 25 6 14 24 36 50 11 24
[13] 39 56 75 16 34 54 76 100 21 44 69 96
[25] 125
```

```
> seq(from = 0, to = 1, by = 0.1) ##A sequence of numbers from 0 to 1
```

```
[1] 0.0 0.1 0.2 0.3 0.4 0.5 0.6 0.7 0.8 0.9 1.0
```

```
> exp(seq(from = 0, to = 1, by = 0.1)) ##What will this produce?
```

```
[1] 1.000 1.105 1.221 1.350 1.492 1.649 1.822 2.014
[9] 2.226 2.460 2.718
```

Now the mysterious square bracketed numbers appearing next to the output make sense. R puts the position of the beginning value on a line in square brackets before the line of output. For example if the output has 40 values, and 15 values appear on each line, then the first line will have [1] at the left, the second line will have [16] to the left, and the third line will have [31] to the left.

2.3.2 Basic descriptive statistics and graphics in R

Of course it is easy to compute basic descriptive statistics and to produce standard graphical representations of data. For illustration consider the first 14 observations of tree height and DBH (diameter at breast height) from the FEF data set. We will begin by entering these data “by hand” using the `c()` function, which concatenates its arguments into a vector. For larger data sets we will clearly want an alternative way to enter data.

A style note: R has two widely used methods of assignment: the left arrow, which consists of a less than sign followed immediately by a dash: `<-` and the equals sign: `=`. Much ink has been used debating the relative merits of the two methods, and their subtle differences. Many leading R style guides (e.g., the Google style guide at <https://google.github.io/styleguide/Rguide.xml> and the Bioconductor style guide at <http://www.bioconductor.org/developers/how-to/coding-style/>) recommend the left arrow `<-` as an assignment operator, and we will use this throughout the book.

Also you will see that if a command has not been completed but the ENTER key is pressed, the command prompt changes to a `+` sign.

```
> dbh <- c(6, 6.9, 6.4, 6.5, 7.2, 3.1, 2, 4.1, 2.4, 2.7, 3.7,  
+ 6.3, 5.2, 5.1, 6.4)  
> ht <- c(48, 48, 48, 49, 51, 40, 30.5, 50, 28, 40.4, 42.6,  
+ 53, 55, 50, 50)  
> dbh
```

```
[1] 6.0 6.9 6.4 6.5 7.2 3.1 2.0 4.1 2.4 2.7 3.7 6.3  
[13] 5.2 5.1 6.4
```

```
> ht
```

```
[1] 48.0 48.0 48.0 49.0 51.0 40.0 30.5 50.0 28.0 40.4  
[11] 42.6 53.0 55.0 50.0 50.0
```

Next we compute some descriptive statistics for the two numeric variables

```
> mean(dbh)
```

```
[1] 4.933
```

```
> sd(dbh)
```

```
[1] 1.782
```

```
> summary(dbh)
```

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
2.00	3.40	5.20	4.93	6.40	7.20

```
> mean(ht)
```

```
[1] 45.57
```

```
> sd(ht)
```


```
[1] 7.857
```

```
> summary(ht)
```

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
28.0	41.5	48.0	45.6	50.0	55.0

Next, a scatter plot of `dbh` versus `ht`:

```
> plot(dbh, ht)
```



Unsurprisingly as DBH increases, height tends to increase. We'll investigate this further using simple linear regression in the next section.

2.3.3 Simple linear regression in R

The `lm()` function is used to fit linear models in R, including simple linear regression models. Here it is applied to the DBH height data.

```
> ht.lm <- lm(ht ~ dbh) ##Fit the model and save it in ht.lm
> summary(ht.lm) ##Basic summary of the model
```

Call:

```
lm(formula = ht ~ dbh)
```

Residuals:

Min	1Q	Median	3Q	Max
-8.507	-2.742	-0.812	2.683	8.480

Coefficients:

Estimate	Std. Error	t value	Pr(> t)

```
(Intercept) 27.925      3.739      7.47 4.7e-06 ***
dbh          3.576      0.716      5.00 0.00024 ***
---
```

Signif. codes:


```
0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

Residual standard error: 4.77 on 13 degrees of freedom

Multiple R-squared: 0.658, Adjusted R-squared: 0.631

F-statistic: 25 on 1 and 13 DF, p-value: 0.000244

```
> plot(dbh, ht) ##Scatter plot of the data
> abline(ht.lm) ##Add the fitted regression line to the plot
```



We will work extensively with such models later in the text. We will also talk about why it might not be a good idea to assume a linear relationship between DBH and height—can you guess why this is by looking at the data scatter and model fitted line in the plot above?

2.4 How to Learn

There are several ways to engage with the content of this book and associated learning materials.

A comprehensive, but slightly overwhelming, cheatsheet for RStudio is available here <https://www.rstudio.com/wp-content/uploads/2016/01/rstudio-IDE-cheatsheet.pdf>. As we progress in learning R and RStudio, this cheatsheet will become more useful. For now you might use the cheatsheet to locate the various windows and functions identified in the coming chapters.

2.5 Getting help

There are several free (and several not free) ways to get R help when needed.

Several help-related functions are built into R. If there's a particular R function of interest, such as `log`, `help(log)` or `?log` will bring up a help page for that function. In RStudio the help page is displayed, by default, in the **Help** tab in the lower right window.⁵ The function `help.start` opens a window which allows browsing of the online documentation included with R. To use this, type `help.start()` in the console window.⁶ The `help.start` function also provides several manuals online and can be a useful interface in addition to the built in help.

Search engines provide another, sometimes more user-friendly, way to receive answers for R questions. A Google search often quickly finds something written by another user who had the same (or a similar) question, or an online tutorial that touches on the question. More specialized is <https://rseek.org/>, which is a search engine focused specifically on R. Both Google and <https://rseek.org> are valuable tools, often providing more user-friendly information than R's own help system.

In addition, R users have written many types of contributed documentation. Some of this documentation is available at <http://cran.r-project.org/other-docs.html>. Of course there are also numerous books covering general and specialized R topics available for purchase.

⁵There are ways to change this default behavior.

⁶You may wonder about the parentheses after `help.start`. A user can specify arguments to any R function inside parentheses. For example `log(10)` asks R to return the logarithm of the argument 10. Even if no arguments are needed, R requires empty parentheses at the end of any function name. In fact if you just type the function name without parentheses, R returns the definition of the function. For simple functions this can be illuminating.

2.6 Workspace, working directory, and keeping organized

The *workspace* is your R session working environment and includes any objects you create. Recall these objects are listed in the **Global Environment** window. The command `ls()`, which stands for list, will also list all the objects in your workspace (note, this is the same list that is given in the **Global Environment** window). When you close RStudio, a dialog box will ask you if you want to save an image of the current workspace. If you choose to save your workspace, RStudio saves your session objects and information in a `.RData` file (the period makes it a hidden file) in your *working directory*. Next time you start R or RStudio it checks if there is a `.RData` in the working directory, loads it if it exists, and your session continues where you left off. Otherwise R starts with an empty workspace. This leads to the next question—what is a working directory?

Each R session is associated with a working directory. This is just a directory from which R reads and writes files, e.g., the `.RData` file, data files you want to analyze, or files you want to save. On Mac when you start RStudio it sets the working directory to your home directory (for me that's `/Users/andy`). If you're on a different operating system, you can check where the default working directory is by typing `getwd()` in the console. You can change the default working directory under RStudio's **Global Option** dialog found under the **Tools** dropdown menu. There are multiple ways to change the working directory once an R session is started in RStudio. One method is to click on the **Files** tab in the lower right window and then click the **More** button. Alternatively, you can set the session's working directory using the `setwd()` in the console. For example, on Windows `setwd("C:/Users/andy/for472/exercise1")` will set the working directory to `C:/Users/andy/for472/exercise1`, assuming that file path and directory exist (Note: Windows file path uses a backslash, `\`, but in R the backslash is an escape character, hence specifying file paths in R on Windows uses the forward slash, i.e., `/`). Similarly on Mac you can use `setwd("/Users/andy/for472/exercise1")`. Perhaps the most simple method is to click on the **Session** tab at the top of your screen and click on the **Set Working Directory** option. Later on when we start reading and writing data from our R session, it will be very important that you are able to identify your current working directory and change it if needed. We will revisit this in subsequent chapters.

As with all work, keeping organized is the key to efficiency. It is good practice to have a dedicated directory for each R project or exercise.

2.7 Quality of R code

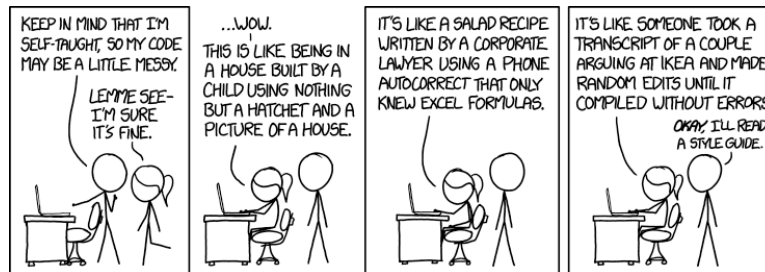


FIGURE 2.2: xkcd: Code Quality

Writing well-organized and well-labeled code allows your code to be more easily read and understood by another person. (See xkcd’s take on code quality in Figure 2.2.) More importantly, though, your well-written code is more accessible to you hours, days, or even months later. We are hoping that you can use the code you write in this class in future projects and research.

Google provides style guides for many programming languages. You can find the R style guide here⁷. Below are a few of the key points from the guide that we will use right away.

2.7.1 Naming Files

File names should be meaningful and end in `.R`. If we write a script that analyzes a certain species distribution:

- GOOD: `african_rhino_distribution.R`
- GOOD: `africanRhinoDistribution.R`
- BAD: `speciesDist.R` (too ambiguous)
- BAD: `species.dist.R` (too ambiguous and two periods can confuse operating systems’ file type auto-detect)
- BAD: `speciesdist.R` (too ambiguous and confusing)

2.7.2 Naming Variables

- GOOD: `rhino.count`
- GOOD: `rhinoCount`

⁷<https://google.github.io/styleguide/Rguide.xml>

- GOOD: `rhino_count` (We don't mind the underscore and use it quite often, although Google's style guide says it's a no-no for some reason)
- BAD: `rhinocount` (confusing)

2.7.3 Syntax

- Keep code lines under 80 characters long.
- Indent your code with two spaces. (RStudio does this by default when you press the TAB key.)

3

Scripts, R Markdown, and Reproducible Research

Doing work in data science, whether for homework, a project for a business, or a research project, typically involves several iterations. For example, creating an effective graphical representation of data can involve trying out several different graphical representations, and then tens if not hundreds of iterations when fine-tuning the chosen representation. And each of these representations may require several R commands to create. Although this all could be accomplished by typing and re-typing commands at the R Console, it is easier and more effective to write the commands in a *script file* that can then be submitted to the R console either a line at a time or all together.¹

In addition to making the workflow more efficient, R scripts provide another large benefit. Often we work on one part of a homework assignment or project for a few hours, then move on to something else, and then return to the original part a few days, months, or sometimes even years later. In such cases we may have forgotten how we created a graphical display that we were so proud of, and will again need to spend a few hours to recreate it. If we save a script file, we have the ingredients immediately available when we return to a portion of a project.²

Next consider a larger scientific endeavor. Ideally a scientific study will be reproducible, meaning that an independent group of researchers (or the original researchers) will be able to duplicate the study. Thinking about data science, this means that all the steps taken when working with the data from a study should be reproducible, from selection of variables to formal data analysis. In principle this can be facilitated by explaining, in words, each step of the work with data. In practice, on the other hand, it is typically difficult or impossible to reproduce a full data analysis based on a written explanation. It is much more effective to include the actual computer code that accomplished the data work in the report, whether the report is a homework assignment or a research paper. Tools in R such as *R Markdown* facilitate this process.

¹Unsurprisingly it is also possible to submit several selected lines of code at once.

²In principle the R history mechanism provides a similar record. But with history we have to search through a lot of other code to find what we're looking for, and scripts are a much cleaner mechanism to record our work.

3.1 Scripts in R

As noted above, scripts help to make working with data more efficient and provide a record of how data were managed and analyzed. Here we describe an example using the FEF data.³ First we read the FEF data into R using the code below.

```
> face.dat <- read.csv(
+   file="http://blue.for.msu.edu/FOR472/data/FACE_aspen_core_growth.csv"
+ )
```

Next we print the names of the variables in the data set. Don't be concerned about the specific details. Later we will learn much more about reading in data and working with data sets in R.

```
> names(face.dat)

[1] "Rep"                "Treat"
[3] "Clone"             "E.Clone"
[5] "Row"               "Col"
[7] "ID.."              "X1997Initial_Height"
[9] "X1997Initial_Diam" "X1997Final_Height"
[11] "X1997Final_Diam"   "X1998_Height"
[13] "X1998_Diam"        "X1999_Height"
[15] "X1999_Diam"        "X2000_Height"
[17] "X2000_Diam"        "X2001_Height"
[19] "X2001_AvgDiam"     "X2001_Diam.3cm"
[21] "X2001_Diam.10cm"   "X2002_Height"
[23] "X2002_Diam.10cm"   "X2003_Height"
[25] "X2003_Diam.10cm"   "X2003_DBH"
[27] "X2004_Height"      "X2004_Diam.10cm"
[29] "X2004_DBH"         "X2005_Height"
[31] "X2005_Diam.10cm"   "X2005_DBH"
[33] "X2006_Height"      "X2006_DBH"
[35] "X2007_Height"      "X2007_DBH"
[37] "X2008_Height"      "X2008_DBH"
[39] "Notes"             "Comment1"
[41] "Comment2"          "Comment3"
[43] "Comment4"          "Comment5"
```

³The example uses features of R that we have not yet discussed, so don't worry about the details but rather about how it motivates the use of a script file.


```
[45] "Comment6"
```

Let's create a scatter plot of 2008 DBH versus height. To do this we'll first create variables for DBH and height taken in the year 2008 and print out the first ten values of each variable.⁴

```
> dbh <- face.dat$X2008_DBH
> ht <- face.dat$X2008_Height
> dbh[1:10]
```


```
[1] NA 9.55 2.00 9.00 3.11 6.35 4.60 NA NA 1.42
```

```
> ht[1:10]
```

```
[1] NA 1225 334 1079 370 859 818 NA NA 268
```

The NA is how missing data are represented in R. Their presence here suggests several trees in this data set are dead or not measured for some reason in 2008. Of course at some point it would be good to investigate which trees have missing data and why. The `plot()` function in R will omit missing values, and for now we will just plot the non-missing data. A scatter plot of the data is drawn next.

```
> plot(dbh, ht)
```




⁴Neither of these steps are necessary, but are convenient for illustration.

Not surprisingly, the scatter plot shows that DBH and height are positively correlated and the relationship is nonlinear. Now that we have a basic scatter plot, it is tempting to make it more informative. We will do this by adding a feature that identifies which trees belong to the control and elevated CO₂ environment treatments. We do this by first separating DBH and height into their respective treatment groups.

```
> treat <- face.dat$Treat
> dbh.treat.1 <- dbh[treat == 1] ##Treatment 1 is the control
> ht.treat.1 <- ht[treat == 1]
>
> dbh.treat.2 <- dbh[treat == 2] ##Treatment 2 is the elevated CO2
> ht.treat.2 <- ht[treat == 2]
```

To make a more informative scatter plot we will do two things. First make a plot for treatment 1 data, but ensure the plot region is large enough to include the treatment 2 data. This is done by specifying the range of the plot axes via `xlim` and `ylim` arguments in the `plot()` function. Here the `xlim` and `ylim` are set to the range of `dbh` and `ht` values, respectively, using `range()` (try and figure out what the `na.rm` argument does in the range function). Second we add treatment 2 data via the `points()` function. There are several other arguments passed to the plot function, but don't worry about these details for now.

```
> plot(dbh.treat.1, ht.treat.1, xlim = range(dbh, na.rm = TRUE),
+      ylim = range(ht, na.rm = TRUE), pch = 19, col = "salmon3",
+      cex = 0.5, xlab = "DBH (cm)", ylab = "Height (cm)")
> points(dbh.treat.2, ht.treat.2, pch = 19, col = "blue",
+        cex = 0.5)
```



Of course we should have a plot legend to tell the viewer which colors are associated with the treatments, as well as many other aesthetic refinements. For now, however, we will resist such temptations.⁵

Some of the process leading to the completed plot is shown above. We read in the data, created an intermediate plot by adding treatment identifiers, creating variables representing the 2008 measurements of DBH and height, and so on. However, a lot of the process isn't shown. For example, I made several mistakes in the process of getting the code and plot the way I wanted it—forgot the `na.rm=TRUE` initially then fiddled around with the treatment colors a bit.

Now imagine trying to recreate the plot a few days later. Possibly someone saw the plot and commented that it would be interesting to see similar plots for each year in the study period. If we did all the work, including all the false starts and refinements, at the console it would be hard to sort things out. This would take much longer than necessary to create the new plots. This would be especially true if a few months had passed, rather than just a few days.

Creating the new scatter plots would be much easier with a script file, especially if it had a few well-chosen comments. Fortunately it is quite easy to create and work with script files in RStudio.⁶ Just choose **File > New File > New script** and a script window will open up in the upper left of the full RStudio window.

An example of a script window (with some R code already typed in) is shown

⁵As an aside, by only looking at the plotted data and thinking about basic plant physiology, can you guess which color is associated with the elevated CO₂ treatment?

⁶It is also easy in R without RStudio. Just use **File > New script** to create a script file, and save it before exiting R.

in Figure 3.1. From the script window the user can, among other things, save the script (either using the **File** menu or the icon near the top left of the window) and can run one or more lines of code from the window (using the **run** icon in the window, or by copying and pasting into the console window). In addition, there is a **Source on Save** checkbox. If this is checked, the R code in the script window is automatically read into R and executed when the script file is saved.

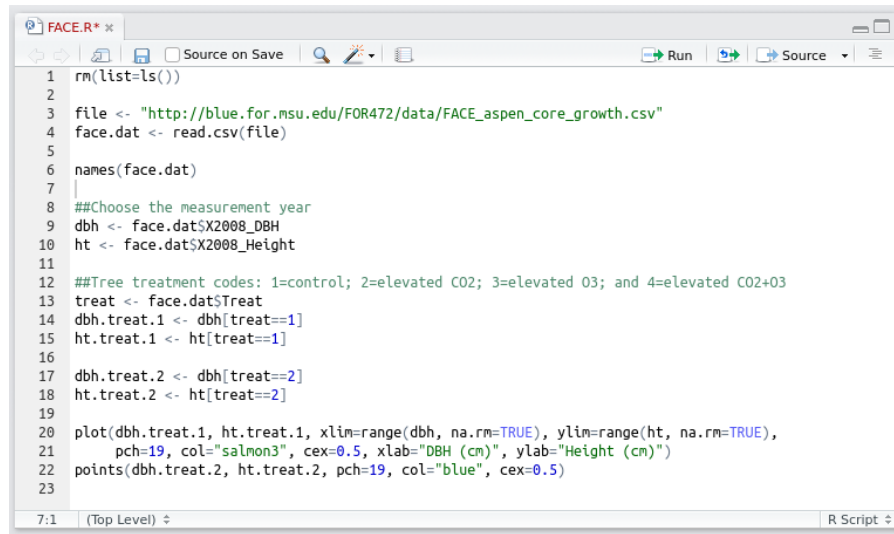


FIGURE 3.1: A script window in RStudio

3.2 R Markdown

People typically work on data with a larger purpose in mind. Possibly the purpose is to understand a biological system more clearly. Possibly the purpose is to refine a system that recommends movies to users in an online streaming movie service. Possibly the purpose is to complete a homework assignment and demonstrate to the instructor an understanding of an aspect of data analysis. Whatever the purpose, a key aspect is communicating with the desired audience.

One possibility, which is somewhat effective, is to write a document using software such as Microsoft Word ⁷ and to include R output such as computations

⁷Or possibly LaTeX if the document is more technical

and graphics by cutting and pasting into the main document. One drawback to this approach is similar to what makes script files so useful: If the document must be revised it may be hard to unearth the R code that created graphics or analyses, to revise these.⁸ A more subtle but possibly more important drawback is that the reader of the document will not know precisely how analyses were done, or how graphics were created. Over time even the author(s) of the paper will forget the details. A verbal description in a “methods” section of a paper can help here, but typically these do not provide all the details of the analysis, but rather might state something like, “All analyses were carried out using R version 3.3.1.”

RStudio’s website provides an excellent overview of R Markdown capabilities for reproducible research. At minimum, follow the **Get Started** link at <http://rmarkdown.rstudio.com/> and watch the introduction video.

Among other things, R Markdown provides a way to include R code that read in data, create graphics, or perform analyses, all in a single document that is processed to create a research paper, homework assignment, or other written product. The R Markdown file is a plain text file containing text the author wants to show in the final document, simple commands to indicate how the text should be formatted (for example boldface, italic, or a bulleted list), and R code that creates output (including graphics) on the fly. Perhaps the simplest way to get started is to see an R Markdown file and the resulting document that is produced after the R Markdown document is processed. In Figure 3.2 we show the input and output of an example R Markdown document. In this case the output created is an HTML file, but there are other possible output formats, such as Microsoft Word or PDF.

At the top of the input R Markdown file are some lines with `---` at the top and the bottom. These lines are not needed, but give a convenient way to specify the title, author, and date of the article that are then typeset prominently at the top of the output document. For now, don’t be concerned with the lines following `output:`. These can be omitted (or included as shown).

Next are a few lines showing some of the ways that font effects such as italics, boldface, and strikethrough can be achieved. For example, an asterisk before and after text sets the text in *italics*, and two asterisks before and after text sets the text in **boldface**.

More important for our purposes is the ability to include R code in the R Markdown file, which will be executed with the output appearing in the output document. Bits of R code included this way are called *code chunks*. The beginning of a code chunk is indicated with three backticks and an “r” in curly

⁸Organizing the R code using script files and keeping all the work organized in a well-thought-out directory structure can help here, but this requires a level of forethought and organization that most people do not possess . . . including myself.

```

---
title: "R Markdown"
author: "Andy Finley"
date: "April 3, 2017"
output: html_document
---

Basic formatting:

*italic*

**bold**

~~strikethrough~~

A code chunk:

```{r}
x <- 1:10
y <- 10:1
mean(x)
sd(y)
```

Inline code:

`r 5+5`

Inline code not executed:

`5+5`

```

R Markdown

Andy Finley

April 3, 2017

Basic formatting:

italic

bold

~~strikethrough~~

A code chunk:

```
x <- 1:10
y <- 10:1
mean(x)
```

```
## [1] 5.5
```

```
sd(y)
```

```
## [1] 3.02765
```

Inline code:

```
10
```

Inline code not executed:

```
5+5
```

FIGURE 3.2: Example R Markdown Input and Output

braces: ````{r}`. The end of a code chunk is indicated with three backticks `````. For example, the R Markdown file in Figure 3.2 has one code chunk:

```

```{r}
x = 1:10
y = 10:1
mean(x)
sd(y)
```

```

In this code chunk two vectors `x` and `y` are created, and the mean of `x` and the standard deviation of `y` are computed. In the output in Figure 3.2 the R code is reproduced, and the output of the two lines of code asking for the mean and standard deviation is shown.

3.2.1 Creating and processing R Markdown documents

RStudio has features which facilitate creating and processing R Markdown documents. Choose **File > New File > R Markdown...** In the ensuing dialog box, make sure that **Document** is highlighted on the left, enter the title and author (if desired), and choose the Default Output Format (HTML is good to begin). Then click OK. A document will appear in the upper left of the RStudio window. It is an R Markdown document, and the title and author you chose will show up, delimited by --- at the top of the document. A generic body of the document will also be included.

For now just keep this generic document as is. To process it to create the HTML output, click the **Knit HTML** button at the top of the R Markdown window⁹. You'll be prompted to choose a filename for the R Markdown file. Make sure that you use `.Rmd` as the extension for this file. Once you've successfully saved the file, RStudio will process the file, create the HTML output, and open this output in a new window. The HTML output file will also be saved to your working directory. This file can be shared with others, who can open it using a web browser such as Chrome or Firefox.

There are many options which allow customization of R Markdown documents. Some of these affect formatting of text in the document, while others affect how R code is evaluated and displayed. The RStudio web site contains a useful summary of many R Markdown options at <https://www.rstudio.com/wp-content/uploads/2015/03/rmarkdown-reference.pdf>. A different, but mind-numbingly busy, cheatsheet is at <https://www.rstudio.com/wp-content/uploads/2016/03/rmarkdown-cheatsheet-2.0.pdf>. Some of the more commonly used R Markdown options are described next.

3.2.2 Text: Lists and Headers

Unordered (sometimes called bulleted) lists and ordered lists are easy in R Markdown. Figure 3.3 illustrates the creation of unordered and ordered lists.

- For an unordered list, either an asterisk, a plus sign, or a minus sign may precede list items. Use a space after these symbols before including the list text. To have second-level items (sub-lists) indent four spaces before indicating the list item. This can also be done for third-level items.
- For an ordered list use a numeral followed by a period and a space (1. or 2. or 3. or ...) to indicate a numbered list, and use a letter followed by a period and a space (a. or b. or c. or ...) to indicate a lettered list. The same four

⁹If you hover your mouse over this Knit button after a couple seconds it should display a keyboard shortcut for you to do this if you don't like pushing buttons

An unordered list:

```
* List item 1
* List item 2
  + Second level list item 1
  + Second level list item 2
    + Third level list item
* List item 3
```

An ordered list:

```
1. List item 1
2. List item 2
  c. Sub list item 1
  q. Sub list item 2
17. List item 3
```

An unordered list:

```
• List item 1
• List item 2
  – Second level list item 1
  – Second level list item 2
    * Third level list item
• List item 3
```

An ordered list:

```
1. List item 1
2. List item 2
  c. Sub list item 1
  d. Sub list item 2
3. List item 3
```

FIGURE 3.3: Producing Lists in R Markdown

space convention used in unordered lists is used to designate ordered sub lists.

- For an ordered list, the first list item will be labeled with the number or letter that you specify, but subsequent list items will be numbered sequentially. The example in Figure 3.3 will make this more clear. In those examples notice that for the ordered list, although the first-level numbers given in the R Markdown file are 1, 2, and 17, the numbers printed in the output are 1, 2, and 3. Similarly the letters given in the R Markdown file are c and q, but the output file prints c and d.

R Markdown does not give substantial control over font size. Different “header” levels are available that provide different font sizes. Put one or more hash marks in front of text to specify different header levels. Other font choices such as subscripts and superscripts are possible, by surrounding the text either by tildes or carets. More sophisticated mathematical displays are also possible, and are surrounded by dollar signs. The actual mathematical expressions are specified using a language called LaTeX See Figures 3.4 and 3.5 for examples.

3.2.3 Code Chunks

R Markdown provides a large number of options to vary the behavior of code chunks. In some contexts it is useful to display the output but not the R code leading to the output. In some contexts it is useful to display the R prompt, while in others it is not. Maybe we want to change the size of figures created

| | |
|--|------------------------------------|
| # A first <i>*level*</i> <i>~~header~~</i> | A first <i>level</i> header |
| ## A second level header | A second level header |
| ### A third level header | A third level header |
| Text subscripts and superscripts: | Text subscripts and superscripts: |
| $x^2 + y^2$ | $x_2 + y_2$ |
| $10^3 = 1000$ | $10^3 = 1000$ |
| Mathematics examples: | Mathematics examples: |
| x_a | x_a |
| x^a | x^a |
| $\int_0^1 x^2 dx$ | $\int_0^1 x^2 dx$ |
| $\frac{x}{y}$ | $\frac{x}{y}$ |
| \sqrt{x} | \sqrt{x} |
| $\sqrt[n]{x}$ | $\sqrt[n]{x}$ |
| $\sum_{k=1}^n$ | $\sum_{k=1}^n$ |
| $\prod_{k=1}^n$ | $\prod_{k=1}^n$ |

FIGURE 3.4: Headers and Some LaTeX in R Markdown

by graphics commands. And so on. A large number of code chunk options are described in <http://www.rstudio.com/wp-content/uploads/2015/03/rmarkdown-reference.pdf>.

Code chunk options are specified in the curly braces near the beginning of a code chunk. Below are a few of the more commonly used options are described. The use of these options is illustrated in Figure 3.6.

1. `echo=FALSE` specifies that the R code itself should not be printed, but any output of the R code should be printed in the resulting document.
2. `include=FALSE` specifies that neither the R code nor the output should be printed. However, the objects created by the code chunk will be available for use in later code chunks.
3. `eval=FALSE` specifies that the R code should not be evaluated. The

| | | | | | | | |
|------------|-----------------------|-----------|----------------------|-----------|----------------------|-------------------|------------------------------|
| \leq | <code>\leq</code> | \geq | <code>\geq</code> | \neq | <code>\neq</code> | \approx | <code>\approx</code> |
| \times | <code>\times</code> | \div | <code>\div</code> | \pm | <code>\pm</code> | \cdot | <code>\cdot</code> |
| $^{\circ}$ | <code>\circ</code> | \circ | <code>\circ</code> | $'$ | <code>\prime</code> | \cdots | <code>\cdots</code> |
| ∞ | <code>\infty</code> | \neg | <code>\neg</code> | \wedge | <code>\wedge</code> | \vee | <code>\vee</code> |
| \supset | <code>\supset</code> | \forall | <code>\forall</code> | \in | <code>\in</code> | \rightarrow | <code>\rightarrow</code> |
| \subset | <code>\subset</code> | \exists | <code>\exists</code> | \notin | <code>\notin</code> | \Rightarrow | <code>\Rightarrow</code> |
| \cup | <code>\cup</code> | \cap | <code>\cap</code> | $ $ | <code>\mid</code> | \Leftrightarrow | <code>\Leftrightarrow</code> |
| \dot{a} | <code>\dot{a}</code> | \hat{a} | <code>\hat{a}</code> | \bar{a} | <code>\bar{a}</code> | \tilde{a} | <code>\tilde{a}</code> |
| α | <code>\alpha</code> | β | <code>\beta</code> | γ | <code>\gamma</code> | δ | <code>\delta</code> |
| ϵ | <code>\epsilon</code> | ζ | <code>\zeta</code> | η | <code>\eta</code> | ε | <code>\varepsilon</code> |
| θ | <code>\theta</code> | ι | <code>\iota</code> | κ | <code>\kappa</code> | ϑ | <code>\vartheta</code> |
| λ | <code>\lambda</code> | μ | <code>\mu</code> | ν | <code>\nu</code> | ξ | <code>\xi</code> |
| π | <code>\pi</code> | ρ | <code>\rho</code> | σ | <code>\sigma</code> | τ | <code>\tau</code> |
| υ | <code>\upsilon</code> | ϕ | <code>\phi</code> | χ | <code>\chi</code> | ψ | <code>\psi</code> |
| ω | <code>\omega</code> | Γ | <code>\Gamma</code> | Δ | <code>\Delta</code> | Θ | <code>\Theta</code> |
| Λ | <code>\Lambda</code> | Ξ | <code>\Xi</code> | Π | <code>\Pi</code> | Σ | <code>\Sigma</code> |
| Υ | <code>\Upsilon</code> | Φ | <code>\Phi</code> | Ψ | <code>\Psi</code> | Ω | <code>\Omega</code> |

FIGURE 3.5: Other useful LaTeX symbols and expressions in R Markdown

code will be printed unless, for example, `echo=FALSE` is also given as an option.

4. `error=FALSE` and `warning=FALSE` specify that, respectively, error messages and warning messages generated by the R code should not be printed.
5. The `comment` option allows a specified character string to be prepended to each line of results. By default this is set to `comment = '##'` which explains the two hash marks preceding the results in Figure 3.2. Setting `comment = NA` presents output without any character string prepended. That is done in most code chunks in this book.
6. `prompt=TRUE` specifies that the R prompt `>` will be prepended to each line of R code shown in the document. `prompt = FALSE` specifies that command prompts should not be included.
7. `fig.height` and `fig.width` specify the height and width of figures generated by R code. These are specified in inches. For example, `fig.height=4` specifies a four inch high figure.

Figures 3.6 gives examples of the use of code chunk options.

```

No options:
```{r}
x <- 1:10
x
```

echo=FALSE:
```{r, echo=FALSE}
x <- 1:10
x
```

comment=NA:
```{r, comment=NA}
x <- 1:10
x
```

comment='#', prompt=TRUE:
```{r, comment='#', prompt=TRUE}
x <- 1:10
x
```

echo=FALSE, fig.height=4, fig.width=4:
```{r, echo=FALSE, fig.height=4, fig.width=4}
y <- 10:1
plot(x,y)
```

No options:
x = 1:10
x

## [1] 1 2 3 4 5 6 7 8 9 10

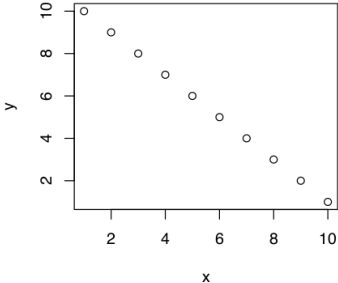
echo=FALSE:
## [1] 1 2 3 4 5 6 7 8 9 10

comment=NA:
x = 1:10
x

[1] 1 2 3 4 5 6 7 8 9 10

comment='#', prompt=TRUE:
> x = 1:10
> x

# [1] 1 2 3 4 5 6 7 8 9 10

echo=FALSE, fig.height=4, fig.width=4:


```

FIGURE 3.6: Output of Example R Markdown

3.2.4 Output formats other than HTML

It is possible to use R Markdown to produce documents in formats other than HTML, including Word and PDF documents. Next to the `Knit HTML` button is a down arrow. Click on this and choose `Knit Word` to produce a Microsoft word output document. Although there is also a `Knit PDF` button, PDF output requires additional software called TeX in addition to RStudio.¹⁰

3.2.5 LaTeX, knitr, and bookdown

While R Markdown provides substantial flexibility and power, it lacks features such as cross-referencing, fine control over fonts, etc. If this is desired, a variant of R Markdown called `knitr`, which has very similar syntax to R Markdown for code chunks, can be used in conjunction with the typesetting system LaTeX to produce documents. We originally created this book using `knitr` and LaTeX. For simpler tasks, however, R Markdown is sufficient, and substantially easier to learn.

As you know (since you are reading this) we are currently converting this book into R Markdown using the package `bookdown` written by Yihui Xie. This package utilizes the R Markdown style that is described above, and also incorporates numerous other features that R Markdown alone does not have (see the previous paragraph). Perhaps the best part about `bookdown` (in addition to its lovely formatting style) is that we can make it interactive, so as you read the html version of this book you can interact with the code itself. You will experience this first hand when you work through the spatial data and databases chapters.

¹⁰It isn't particularly hard to install TeX software. For a Microsoft Windows system, MiKTeX is convenient and is available from <https://miktex.org>. For a Mac system, MacTeX is available from <https://www.tug.org/mactex/>

4

Data Structures

A data structure is a format for organizing and storing data. The structure is designed so that data can be accessed and worked with in specific ways. Statistical software and programming languages have methods (or functions) designed to operate on different kinds of data structures.

This chapter's focus is on data structures. To help initial understanding, the data in this chapter will be relatively modest in size and complexity. The ideas and methods, however, generalize to larger and more complex data sets.

The base data structures in R are vectors, matrices, arrays, data frames, and lists. The first three, vectors, matrices, and arrays, are *homogeneous*, meaning that all elements are required to be of the same type (e.g., all numeric or all character). Data frames and lists are *heterogeneous*, allowing elements to be of different types (e.g., some elements of a data frame may be numeric while other elements may be character). These base structures can also be organized by their dimensionality, as shown in Table 4.1.

R has no scalar types (0-dimensional). Individual numbers or strings are actually vectors of length one.

An efficient way to understand what comprises a given object is to use the `str()` function. `str()` is short for structure and prints a compact, human-readable description of any R data structure. For example, in the code below, we prove to ourselves that what we might think of as a scalar value is actually a vector of length one.

```
> a <- 1
> str(a)
```

TABLE 4.1: Dimension and Type Content of Base Data Structures in R.

| Dimension | Homogeneous | Heterogeneous |
|-----------|---------------|---------------|
| 1 | Atomic Vector | List |
| 2 | Matrix | Data Frame |
| N | Array | |

```
num 1
```

```
> is.vector(a)
```

```
[1] TRUE
```

```
> length(a)
```

```
[1] 1
```

Here we assigned `a` the scalar value one. The `str(a)` prints `num 1`, which says `a` is numeric of length one. Then just to be sure we used the function `is.vector()` to test if `a` is in fact a vector. Then, just for fun, we computed the length of `a` which again returns one. There are a set of similar logical tests for the other base data structures, e.g., `is.matrix()`, `is.array()`, `is.data.frame()`, and `is.list()`. These will all come in handy as we encounter different R objects.

4.1 Vectors

Think of a vector¹ as a structure to represent one variable in a data set. For example a vector might hold the DBH, in inches, of six trees in a data set, and another vector might hold the species of those six trees. The `c()` function in R is useful for creating vectors and for modifying existing vectors. Think of `c` as standing for “combine” or “concatenate.”

```
> dbh <- c(20, 18, 13, 16, 10, 14)
> dbh
```

```
[1] 20 18 13 16 10 14
```

```
> spp <- c("Acer rubrum", "Acer rubrum", "Betula lenta", "Betula lenta",
+         "Prunus serotina", "Prunus serotina")
> spp
```

```
[1] "Acer rubrum"      "Acer rubrum"
```

¹Technically the objects described in this section are “atomic” vectors (all elements of the same type), since lists are also actually vectors. This will not be an important issue in this course, and the shorter term vector will be used for atomic vectors.

```
[3] "Betula lenta"      "Betula lenta"  
[5] "Prunus serotina"   "Prunus serotina"
```

Notice that elements of a vector are separated by commas when using the `c()` function to create a vector. Also notice that character values are placed inside quotation marks.

The `c()` function also can be used to add to an existing vector. For example, if a seventh tree were included in the data set, and its DBH was 13 inches, the existing vectors could be modified as follows.

```
> dbh <- c(dbh, 13)  
> spp <- c(spp, "Acer rubrum")  
> dbh
```

```
[1] 20 18 13 16 10 14 13
```

```
> spp
```

```
[1] "Acer rubrum"      "Acer rubrum"  
[3] "Betula lenta"     "Betula lenta"  
[5] "Prunus serotina" "Prunus serotina"  
[7] "Acer rubrum"
```

4.1.1 Types, Conversion, and Coercion

Clearly it is important to distinguish between different types of vectors. For example, it makes sense to ask R to calculate the mean of the DBH stored in `dbh`, but does not make sense to ask R to compute the mean of the species stored in `spp`. Vectors in R may have one of six different “types”: character, double, integer, logical, complex, and raw. Only the first four of these will be of interest below, and the distinction between double and integer will not be of great import. To illustrate logical vectors, imagine the field technician who measured the trees also indicated if the tree was acceptable growing stock (`ags`) and the call was coded as `TRUE` if the tree was acceptable and `FALSE` if the tree was not acceptable.

```
> typeof(dbh)
```

```
[1] "double"
```

```
> typeof(spp)
```

```
[1] "character"
```

```
> ags <- c(TRUE, TRUE, FALSE, TRUE, FALSE, FALSE, TRUE)
> ags
```

```
[1] TRUE TRUE FALSE TRUE FALSE FALSE TRUE
```

```
> typeof(ags)
```

```
[1] "logical"
```

It may be surprising to see the DBH variable `dbh` is of type `double`, even though its values are all integers. By default R creates a double type vector when numeric values are given via the `c()` function.

When it makes sense, it is possible to convert vectors to a different type. Consider the following examples.

```
> dbh.int <- as.integer(dbh)
> dbh.int
```

```
[1] 20 18 13 16 10 14 13
```

```
> typeof(dbh.int)
```

```
[1] "integer"
```

```
> dbh.char <- as.character(dbh)
> dbh.char
```

```
[1] "20" "18" "13" "16" "10" "14" "13"
```

```
> ags.double <- as.double(ags)
> ags.double
```

```
[1] 1 1 0 1 0 0 1
```

```
> spp.oops <- as.double(spp)
```

Warning: NAs introduced by coercion


```
> spp.oops
```

```
[1] NA NA NA NA NA NA NA
```

```
> sum(ags)
```

```
[1] 4
```

The integer version of `dbh` doesn't look any different, but it is stored differently, which can be important both for computational efficiency and for interfacing with other languages such as C++. As noted above, however, we will not worry about the distinction between integer and double types. Converting `dbh` to character goes as expected—the character representation of the numbers replace the numbers themselves. Converting the logical vector `ags` to double is pretty straightforward too—`FALSE` is converted to zero, and `TRUE` is converted to one. Now think about converting the character vector `spp` to a numeric double vector. It's not at all clear how to represent “Acer rubrum” as a number. In fact in this case what R does is to create a double vector, but with each element set to `NA`, which is the representation of missing data². Finally consider the code `sum(ags)`. Now `ags` is a logical vector, but when R sees that we are asking to sum this logical vector, it automatically converts it to a numerical vector and then adds the zeros and ones representing `FALSE` and `TRUE`.

R also has functions to test whether a vector is of a particular type.

```
> is.double(dbh)
```

```
[1] TRUE
```

```
> is.character(dbh)
```

```
[1] FALSE
```

```
> is.integer(dbh.int)
```

```
[1] TRUE
```

```
> is.logical(ags)
```

```
[1] TRUE
```

²Missing data will be discussed in more detail later in the chapter.

4.1.1.1 Coercion

Consider the following examples.

```
> xx <- c(1, 2, 3, TRUE)
> xx
```

```
[1] 1 2 3 1
```

```
> yy <- c(1, 2, 3, "dog")
> yy
```

```
[1] "1" "2" "3" "dog"
```

```
> zz <- c(TRUE, FALSE, "cat")
> zz
```

```
[1] "TRUE" "FALSE" "cat"
```

```
> dbh + ags
```

```
[1] 21 19 13 17 10 14 14
```

Vectors in R can only contain elements of one type. If more than one type is included in a `c()` function, R silently *coerces* the vector to be of one type. The examples illustrate the hierarchy—if any element is a character, then the whole vector is character. If some elements are numeric (either integer or double) and other elements are logical, then the whole vector is numeric. Note what happened when R was asked to add the numeric vector `dbh` to the logical vector `ags`. The logical vector was silently coerced to be numeric, so that `FALSE` became zero and `TRUE` became one, and then the two numeric vectors were added.

4.1.2 Accessing Specific Elements of Vectors

To access and possibly change specific elements of vectors, refer to the position of the element in square brackets. For example, `dbh[4]` refers to the fourth element of the vector `dbh`. Note that R starts the numbering of elements at 1, i.e., the first element of a vector `x` is `x[1]`.

```
> dbh
```

```
[1] 20 18 13 16 10 14 13
```

```
> dbh[5]
```

```
[1] 10
```

```
> dbh[1:3]
```

```
[1] 20 18 13
```

```
> length(dbh)
```

```
[1] 7
```

```
> dbh[length(dbh)]
```

```
[1] 13
```

```
> dbh[]
```

```
[1] 20 18 13 16 10 14 13
```

```
> dbh[3] <- 202 ##crazy big tree  
> dbh
```

```
[1] 20 18 202 16 10 14 13
```

```
> dbh[1:3] <- c(16, 8, 2)  
> dbh
```

```
[1] 16 8 2 16 10 14 13
```

Note that including nothing in the square brackets results in the whole vector being returned. We can also assign values to vectors by accessing the position(s) where the new values will be assigned. For example, in the above code chunk `dbh[3]` is changed to 202, then the values in the first three elements of `dbh` are changed to 10, 8, and 2, respectively.

Negative numbers in the square brackets tell R to omit the corresponding value. And a zero as a subscript returns nothing (more precisely, it returns a length zero vector of the appropriate type).

```
> dbh[-3]
```

```
[1] 16  8 16 10 14 13
```

```
> dbh[-length(dbh)]
```

```
[1] 16  8  2 16 10 14
```

```
> fewer.dbh <- dbh[-c(1, 3, 5)]  
> fewer.dbh
```

```
[1]  8 16 14 13
```

```
> dbh[0]
```

```
numeric(0)
```

```
> dbh[c(0, 2, 1)]
```

```
[1]  8 16
```

```
> dbh[c(-1, 2)]
```

Error in dbh[c(-1, 2)]: only 0's may be mixed with negative subscripts

Note that mixing zero and other nonzero subscripts is allowed, but mixing negative and positive subscripts is not allowed.

What about the (usual) case where we don't know the positions of the elements we want? For example possibly we want the DBH of all acceptable growing stock trees in the data. Later we will learn how to subset using logical indices, which is a very powerful way to access desired elements of a vector ³.

4.2 Factors

Categorical variables such as `spp` can be represented as character vectors. In many cases this simple representation is sufficient. Consider, however, two

³We had a prelude to this in the temperature data exercise.

other categorical variables, one representing crown class **S**, **I**, **C**, and **D** (i.e., **S**uppressed, **I**ntermediate, **C**odominant, **D**ominant), and another representing grade of the first log via categories **Grade 1**, **Grade 2**, and **Grade 3**. Suppose that for the small data set considered here, all trees are either dominant or codominant crown class. If we just represented the variable via a character vector, there would be no way to know that there are two other categories, representing suppressed and intermediate, because they happen to not be present in the data set. In addition, for the log grade variable the character vector representation does not explicitly indicate that there is an ordering of the levels.

Factors in R provide a more sophisticated way to represent categorical variables. Factors explicitly contain all possible levels, and allow ordering of levels.

```
> crown.class <- c("D", "D", "C", "D", "D", "C", "C")
> grade <- c("Grade 1", "Grade 1", "Grade 3", "Grade 2", "Grade 3",
+ "Grade 3", "Grade 2")
> crown.class
```

```
[1] "D" "D" "C" "D" "D" "C" "C"
```

```
> grade
```

```
[1] "Grade 1" "Grade 1" "Grade 3" "Grade 2" "Grade 3"
[6] "Grade 3" "Grade 2"
```

```
> crown.class <- factor(crown.class, levels = c("S", "I",
+ "C", "D"))
> crown.class
```

```
[1] D D C D D C C
Levels: S I C D
```

```
> grade <- factor(grade, levels = c("Grade 1", "Grade 2",
+ "Grade 3"), ordered = TRUE)
> grade
```

```
[1] Grade 1 Grade 1 Grade 3 Grade 2 Grade 3 Grade 3
[7] Grade 2
Levels: Grade 1 < Grade 2 < Grade 3
```

In the factor version of crown class the levels are explicitly listed, so it is clear that the two included levels are not all the possible levels. In the factor version of log grade, the ordering is explicit as well.

In many cases the character vector representation of a categorical variable is sufficient and easier to work with. In this book factors will not be used extensively. It is important to note that R often by default creates a factor when character data are read in, and sometimes it is necessary to use the argument `stringsAsFactors = FALSE` to explicitly tell R not to do this. This will be seen below when data frames are introduced.

4.3 Names of objects in R

Continuing the discussion about code quality from Section 2.7, there are a few hard and fast restrictions on the names of objects (such as vectors) in R. Note also that there are good practices, and things to avoid.

From the help page for `make.names` in R, the name of an R object is “syntactically valid” if the name “consists of letters, numbers and the dot or underline characters and starts with a letter or the dot not followed by a number” and is not one of the “reserved words” in R such as `if`, `TRUE`, `function`, etc. For example, `c45t.le_dog` and `.ty56` are both syntactically valid (although not very good names) while `4DislikeCats` and `log#@sparty` are not.

A few important comments about naming objects follow:

1. It is important to be aware that names of objects in R are case-sensitive, so `dbh` and `DBH` do not refer to the same object.

```
> dbh
```

```
[1] 16  8  2 16 10 14 13
```

```
> DBH
```

```
Error in eval(expr, envir, enclos): object 'DBH' not found
```

2. It is unwise to create an object with the same name as a built-in R object such as the function `c` or the function `mean`. In earlier versions of R this could be somewhat disastrous, but even in current versions, it is definitely not a good idea!
3. As much as possible, choose names that are informative. When creating a variable you may initially remember that `x` contains `DBH` and `y` contains crown class, but after a few hours, days, or weeks, you probably will forget this. Better options are `dbh` and `crown.class`.

4. As much as possible, be consistent in how you name objects. In particular, decide how to separate multi-word names. Some options include:
 - Using case to separate: `CrownClass` or `crownClass` for example
 - Using underscores to separate: `crown_class` for example
 - Using a period to separate: `crown.class` for example

4.4 Missing Data, Infinity, etc.

Most real-world data sets have variables where some observations are missing. In longitudinal studies of tree growth (i.e., where trees are measured over time), it is common that trees die or cannot be located in subsequent remeasurements. Statistical software should be able to represent missing data and to analyze data sets in which some data are missing.

In R, the value `NA` is used for a missing data value. Since missing values may occur in numeric, character, and other types of data, and since R requires that a vector contain only elements of one type, there are different types of `NA` values. Usually R determines the appropriate type of `NA` value automatically. It is worth noting the default type for `NA` is logical, and that `NA` is NOT the same as the character string `"NA"`.

```
> missingCharacter <- c("dog", "cat", NA, "pig", NA, "horse")
> missingCharacter
```

```
[1] "dog"  "cat"  NA      "pig"  NA      "horse"
```

```
> is.na(missingCharacter)
```

```
[1] FALSE FALSE  TRUE FALSE  TRUE FALSE
```

```
> missingCharacter <- c(missingCharacter, "NA")
> missingCharacter
```

```
[1] "dog"  "cat"  NA      "pig"  NA      "horse"
[7] "NA"
```

```
> is.na(missingCharacter)
```

```
[1] FALSE FALSE  TRUE FALSE  TRUE FALSE FALSE
```

```
> allMissing <- c(NA, NA, NA)
> typeof(allMissing)
```

```
[1] "logical"
```

How should missing data be treated in computations, such as finding the mean or standard deviation of a variable? One possibility is to return `NA`. Another is to remove the missing value(s) and then perform the computation.

```
> mean(c(1, 2, 3, NA, 5))
```

```
[1] NA
```

```
> mean(c(1, 2, 3, NA, 5), na.rm = TRUE)
```

```
[1] 2.75
```

As this example shows, the default behavior for the `mean()` function is to return `NA`. If removal of the missing values and then computing the mean is desired, the argument `na.rm` is set to `TRUE`. Different R functions have different default behaviors, and there are other possible actions. Consulting the help for a function provides the details.

4.4.1 Infinity and NaN

What happens if R code requests division by zero, or results in a number that is too large to be represented? Here are some examples.

```
> x <- 0:4
> x
```

```
[1] 0 1 2 3 4
```

```
> 1/x
```

```
[1]      Inf 1.0000 0.5000 0.3333 0.2500
```



```
> x/x
```

```
[1] NaN  1  1  1  1
```

```
> y <- c(10, 1000, 10000)
> 2^y
```

```
[1] 1.024e+03 1.072e+301      Inf
```

`Inf` and `-Inf` represent infinity and negative infinity (and numbers which are too large in magnitude to be represented as floating point numbers). `NaN` represents the result of a calculation where the result is undefined, such as dividing zero by zero. All of these are common to a variety of programming languages, including R.

4.5 Data Frames

Commonly data is rectangular in form, with variables as columns and cases as rows. Continuing with the species, DBH, and acceptable growing stock data, each of those variables would be a column of the data set, and each tree's measurements would be a row. In R, such data are represented as a *data frame*.

```
> trees <- data.frame(Spp=spp, Dbh=dbh, Ags=ags,
+                      stringsAsFactors=FALSE)
> trees
```

| | Spp | Dbh | Ags |
|---|-----------------|-----|-------|
| 1 | Acer rubrum | 16 | TRUE |
| 2 | Acer rubrum | 8 | TRUE |
| 3 | Betula lenta | 2 | FALSE |
| 4 | Betula lenta | 16 | TRUE |
| 5 | Prunus serotina | 10 | FALSE |
| 6 | Prunus serotina | 14 | FALSE |
| 7 | Acer rubrum | 13 | TRUE |

```
> names(trees)
```

```
[1] "Spp" "Dbh" "Ags"
```

```
> colnames(trees)
```

```
[1] "Spp" "Dbh" "Ags"
```

```
> names(trees) <- c("species", "DBH", "good.stock")
> colnames(trees)
```

```
[1] "species"      "DBH"          "good.stock"
```

```
> rownames(trees)
```

```
[1] "1" "2" "3" "4" "5" "6" "7"
```

```
> names(trees) <- c("spp", "dbh", "ags")
> dim(trees)
```

```
[1] 7 3
```

The `data.frame` function can be used to create a data frame (although it's more common to read a data frame into R from an external file, something that will be introduced later). The names of the variables in the data frame are given as arguments, as are the vectors of data that make up the variable's values. The argument `stringsAsFactors=FALSE` asks R not to convert character vectors into factors, which R does by default, to the dismay of many users. Names of the columns (variables) can be extracted and set via either `names` or `colnames`. In the example, the variable names are changed to `species`, `DBH`, `good.stock` and then changed back to what I like better `spp`, `dbh`, `ags` in this way. Rows can be named also. In this case since specific row names were not provided the default row names of "1", "2", etc. are used. Finally, I take a look at the data frame's dimensions (where the `dim` function returns a vector comprised of number of rows and number of columns, respectively). Also, try the functions `nrow` and `ncol` on the data frame and see what happens.

In the next example a built-in R data set called `Loblolly` is made available by the `data` function, and then the first and last six rows are displayed using `head` and `tail`.

```
> data("Loblolly")
> head(Loblolly)
```

```
      height age Seed
1       4.51  3  301
15    10.89  5  301
```

```
29 28.72 10 301
43 41.74 15 301
57 52.70 20 301
71 60.92 25 301
```

```
> tail(Loblolly)
```

```
      height age Seed
14    3.46   3  331
28    9.05   5  331
42   25.85  10  331
56   39.15  15  331
70   49.12  20  331
84   59.49  25  331
```

Note the `Loblolly` data frame has row names that are not ordered (which really doesn't matter) and simply suggests the data set author might have subset these data from a larger data set or sorted them by a variable, e.g., `height` or `age`. Row names can be generally ignored (unless they hold some specific meaning). Find out more about the `Loblolly` data set by running `?Loblolly` on the command line or, equivalently, looking it up in the RStudio's search window on the help tab.

4.5.1 Accessing specific elements of data frames

Data frames are two-dimensional, so to access a specific element (or elements) we need to specify both the row and column indices.

```
> Loblolly[1, 3]
```

```
[1] 301
14 Levels: 329 < 327 < 325 < 307 < 331 < ... < 305
```

```
> Loblolly[1:3, 3]
```

```
[1] 301 301 301
14 Levels: 329 < 327 < 325 < 307 < 331 < ... < 305
```

```
> Loblolly[1:3, 2:3]
```

```
      age Seed
1      3  301
```

```
15  5 301
29 10 301
```

```
> Loblolly[, 1]
```

```
[1] 4.51 10.89 28.72 41.74 52.70 60.92 4.55 10.92
[9] 29.07 42.83 53.88 63.39 4.79 11.37 30.21 44.40
[17] 55.82 64.10 3.91 9.48 25.66 39.07 50.78 59.07
[25] 4.81 11.20 28.66 41.66 53.31 63.05 3.88 9.40
[33] 25.99 39.55 51.46 59.64 4.32 10.43 27.16 40.85
[41] 51.33 60.07 4.57 10.57 27.90 41.13 52.43 60.69
[49] 3.77 9.03 25.45 38.98 49.76 60.28 4.33 10.79
[57] 28.97 42.44 53.17 61.62 4.38 10.48 27.93 40.20
[65] 50.06 58.49 4.12 9.92 26.54 37.82 48.43 56.81
[73] 3.93 9.34 26.08 37.79 48.31 56.43 3.46 9.05
[81] 25.85 39.15 49.12 59.49
```

Note that `Loblolly[,1]` returns ALL elements in the first column. This agrees with the behavior for vectors, where leaving a subscript out of the square brackets tells R to return all values. In this case we are telling R to return all rows, and the first column.

As we have seen in class, we can also access the columns (or rows) using their names.

```
> Loblolly[1:4, "height"]
```

```
[1] 4.51 10.89 28.72 41.74
```

```
> Loblolly[1:4, c("age", "Seed")]
```

```
      age Seed
1       3 301
15      5 301
29     10 301
43     15 301
```

For a data frame there is another way to access specific columns, using the `$` notation.

```
> Loblolly$height
```

```
[1] 4.51 10.89 28.72 41.74 52.70 60.92 4.55 10.92
[9] 29.07 42.83 53.88 63.39 4.79 11.37 30.21 44.40
```

```
[17] 55.82 64.10  3.91  9.48 25.66 39.07 50.78 59.07
[25]  4.81 11.20 28.66 41.66 53.31 63.05  3.88  9.40
[33] 25.99 39.55 51.46 59.64  4.32 10.43 27.16 40.85
[41] 51.33 60.07  4.57 10.57 27.90 41.13 52.43 60.69
[49]  3.77  9.03 25.45 38.98 49.76 60.28  4.33 10.79
[57] 28.97 42.44 53.17 61.62  4.38 10.48 27.93 40.20
[65] 50.06 58.49  4.12  9.92 26.54 37.82 48.43 56.81
[73]  3.93  9.34 26.08 37.79 48.31 56.43  3.46  9.05
[81] 25.85 39.15 49.12 59.49
```

```
> Loblolly$age
```

```
[1]  3  5 10 15 20 25  3  5 10 15 20 25  3  5 10 15 20
[18] 25  3  5 10 15 20 25  3  5 10 15 20 25  3  5 10 15
[35] 20 25  3  5 10 15 20 25  3  5 10 15 20 25  3  5 10
[52] 15 20 25  3  5 10 15 20 25  3  5 10 15 20 25  3  5
[69] 10 15 20 25  3  5 10 15 20 25  3  5 10 15 20 25
```

```
> height
```

```
Error in eval(expr, envir, enclos): object 'height' not found
```

```
> age
```

```
Error in eval(expr, envir, enclos): object 'age' not found
```

Notice that typing the variable name, such as `height`, without the name of the data frame (and a dollar sign) as a prefix, does not work. This is sensible. There may be several data frames that have variables named `height`, and just typing `height` doesn't provide enough information to know which is desired.

4.6 Lists

The third main data structure we will work with is a list. Technically a list is a vector, but one in which elements can be of different types. For example a list may have one element that is a vector, one element that is a data frame, and another element that is a function. Consider designing a function that fits a simple linear regression model to two quantitative variables. We might want that function to compute and return several things such as

- The fitted slope and intercept (a numeric vector with two components)

- The residuals (a numeric vector with n components, where n is the number of data points)
- Fitted values for the data (a numeric vector with n components, where n is the number of data points)
- The names of the dependent and independent variables (a character vector with two components)

In fact R has a function, `lm`, which does this (and much more).

```
> htAgeLinMod <- lm(height ~ age, data = Loblolly)
> mode(htAgeLinMod)
```

```
[1] "list"
```

```
> names(htAgeLinMod)
```

```
[1] "coefficients" "residuals"      "effects"
[4] "rank"         "fitted.values"  "assign"
[7] "qr"           "df.residual"    "xlevels"
[10] "call"         "terms"         "model"
```

```
> htAgeLinMod$coefficients
```

```
(Intercept)      age
      -1.312      2.591
```

```
> tail(htAgeLinMod$residuals)
```

```
      14      28      42      56      70      84
-2.999 -2.590  1.257  1.605 -1.378 -3.961
```

The `lm` function returns a list (which in the code above has been assigned to the object `htAgeLinMod`)⁴. One component of the list is the length 2 vector of coefficients, while another component is the length 84 vector of residuals. The code also illustrates that named components of a list can be accessed using the dollar sign notation, as with data frames.

The `list` function is used to create lists.

```
> temporaryList <- list(first = dbh, second = trees, pickle = list(a = 1:10,
+   b = trees))
> temporaryList
```

⁴The `mode` function returns the type or storage mode of an object.

```

$first
[1] 16  8  2 16 10 14 13

$second
      spp dbh  ags
1   Acer rubrum 16 TRUE
2   Acer rubrum  8 TRUE
3  Betula lenta  2 FALSE
4  Betula lenta 16 TRUE
5 Prunus serotina 10 FALSE
6 Prunus serotina 14 FALSE
7   Acer rubrum 13 TRUE

$pickle
$pickle$a
[1] 1 2 3 4 5 6 7 8 9 10

$pickle$b
      spp dbh  ags
1   Acer rubrum 16 TRUE
2   Acer rubrum  8 TRUE
3  Betula lenta  2 FALSE
4  Betula lenta 16 TRUE
5 Prunus serotina 10 FALSE
6 Prunus serotina 14 FALSE
7   Acer rubrum 13 TRUE

```

Here, for illustration, I assembled a list to hold some of the R data structures we have been working with in this chapter. The first list element, named `first`, holds the `dbh` vector we created in Section 4.1. The second list element, named `second`, holds the `trees` data frame. The third list element, named `pickle`, holds a list with elements named `a` and `b` that hold a vector of values 1 through 10, and another copy of the `trees` data set, respectively. As this example shows, a list can contain another list.

4.6.1 Accessing specific elements of lists

We already have seen the dollar sign notation works for lists. In addition, the square bracket subsetting notation can be used. But with lists there is an added somewhat subtle wrinkle—using either single or double square brackets.

```
> temporaryList$first
```

```
[1] 16  8  2 16 10 14 13
```

```
> mode(temporaryList$first)
```

```
[1] "numeric"
```

```
> temporaryList[[1]]
```

```
[1] 16  8  2 16 10 14 13
```

```
> mode(temporaryList[[1]])
```

```
[1] "numeric"
```

```
> temporaryList[1]
```

```
$first
[1] 16  8  2 16 10 14 13
```

```
> mode(temporaryList[1])
```

```
[1] "list"
```

Note the dollar sign and double bracket notation return a numeric vector, while the single bracket notation returns a list. Notice also the difference in results below.

```
> temporaryList[c(1, 2)]
```

```
$first
[1] 16  8  2 16 10 14 13
```

```
$second
      spp dbh  ags
1   Acer rubrum 16 TRUE
2   Acer rubrum  8 TRUE
3  Betula lenta  2 FALSE
4  Betula lenta 16 TRUE
5 Prunus serotina 10 FALSE
6 Prunus serotina 14 FALSE
7   Acer rubrum 13 TRUE
```



```
> temporaryList[[c(1, 2)]]
```

```
[1] 8
```

The single bracket form returns the first and second elements of the list, while the double bracket form returns the second element in the first element of the list. Generally, do not put a vector of indices or names in a double bracket, you will likely get unexpected results. See, for example, the results below ⁵.

```
> temporaryList[[c(1, 2, 3)]]
```

```
Error in temporaryList[[c(1, 2, 3)]]: recursive indexing failed at level 2
```

So, in summary, there are two main differences between using the single bracket `[]` and double bracket `[[]]`. First, the single bracket will return a list that holds the object(s) held at the given indices or names placed in the bracket, whereas the double brackets will return the actual object held at the index or name placed in the innermost bracket. Put differently, a single bracket can be used to access a range of list elements and will return a list, while a double bracket can only access a single element in the list and will return the object held at the index.

4.7 Subsetting with Logical Vectors

Consider the `Loblolly` data frame. How can we access only those trees with heights more than 50 m? How can we access the age of those trees taller than 50 m? How can we compute the mean height of all trees from seed source 301? The data set is small enough that it would not be too onerous to extract the values by hand. But for larger or more complex data sets, this would be very difficult or impossible to do in a reasonable amount of time, and would likely result in errors.

R has a powerful method for solving these sorts of problems using a variant of the subsetting methods that we already have learned. When given a logical vector in square brackets, R will return the values corresponding to `TRUE`. To begin, focus on the `dbh` and `spp` vectors created in Section 4.1.

The R code `dbh > 15` returns `TRUE` for each value of `dbh` that is more than 15, and a `FALSE` for each value of `dbh` that is less than or equal to 15. Similarly `spp`

⁵Try this example using only single brackets... it will return a list holding elements `first`, `second`, and `pickle`.

`== "Betula lenta"` returns `TRUE` or `FALSE` depending on whether an element of `spp` is equal to `Betula lenta`.

```
> dbh
```

```
[1] 16  8  2 16 10 14 13
```

```
> dbh > 15
```

```
[1]  TRUE FALSE FALSE  TRUE FALSE FALSE FALSE
```

```
> spp[dbh > 15]
```

```
[1] "Acer rubrum" "Betula lenta"
```

```
> dbh[dbh > 15]
```

```
[1] 16 16
```

```
> spp == "Betula lenta"
```

```
[1] FALSE FALSE  TRUE  TRUE FALSE FALSE FALSE
```

```
> dbh[spp == "Betula lenta"]
```

```
[1]  2 16
```

Consider the lines of R code one by one.

- `dbh` instructs R to display the values in the vector `dbh`.
- `dbh > 15` instructs R to check whether each value in `dbh` is greater than 15, and to return `TRUE` if so, and `FALSE` otherwise.
- The next line, `spp[dbh > 15]`, does two things. First, inside the square brackets, it does the same thing as the second line: it returns `TRUE` or `FALSE` depending on whether a value of `dbh` is or is not greater than 15. Second, each element of `spp` is matched with the corresponding `TRUE` or `FALSE` value, and is returned if and only if the corresponding value is `TRUE`. For example the first value of `spp` is `Acer rubrum`. Since the first `TRUE` or `FALSE` value is `TRUE`, the first value of `spp` is returned. So this line of code returns the species names for all trees with DBH greater than 15; hence, the first and the fourth values of `spp` are returned.
- The fourth line of code, `dbh[dbh > 15]`, again begins by returning `TRUE` or `FALSE` depending on whether elements of `dbh` are larger than 15. Then those

elements of `dbh` corresponding to `TRUE` values are returned. So this line of code returns the DBH of all trees whose DBH is greater than 15.

- The fifth line returns `TRUE` or `FALSE` depending on whether elements of `spp` are equal to `Betula lenta` or not.
- The sixth line returns the DBH of those all `Betula lenta` trees.

There are six comparison operators in R, `>`, `<`, `>=`, `<=`, `==`, `!=`. Note that to test for equality a “double equals sign” is used, while `!=` tests for inequality.



Bibliography

- Asner, G., Hughes, R., Varga, T., Knapp, D., and Kennedy-Bowdoin, T. (2009). Environmental and biotic controls over aboveground biomass throughout a tropical rain forest. *Ecosystems*, 12(2):261–278.
- Babcock, C., Matney, J., Finley, A., Weiskittel, A., and Cook, B. (2013). Multivariate spatial regression models for predicting individual tree structure variables using lidar data. *IEEE Journal of Selected Topics in Applied Earth Observations and Remote Sensing*, 6(1, SI):6–14.
- Bravo, A., Porzecanski, A., Sterling, E., Bynum, N., Cawthorn, M., Fernandez, D. S., Freeman, L., Ketcham, S., Leslie, T., Mull, J., and Vogler, D. (2016). Teaching for higher levels of thinking: developing quantitative and analytical skills in environmental science courses. *Ecosphere*, 7(4):n/a–n/a. e01290.
- Finley, A., Banerjee, S., and MacFarlane, D. (2011). A hierarchical model for quantifying forest variables over large heterogeneous landscapes with uncertain forest areas. *Journal of the American Statistical Association*, 106(493):31–48.
- Kenefic, L. S., Rogers, N. S., Puhlick, J. J., Waskiewicz, J. D., and Brissette, J. C. (2015). Overstory tree and regeneration data from the "silvicultural effects on composition, structure, and growth" study at penobscot experimental forest. 2nd edition. fort collins, co: Forest service research data archive.
- Kubiske, M. E. (2013). Tree height and diameter data from the aspen face experiment, 1997-2008. newtown square, pa: Usda forest service, northern research station.
- Le Toan, T., Quegan, S., Davidson, M., Balzter, H., Paillou, P., Papathanassiou, K., Plummer, S., Rocca, F., Saatchi, S., Shugart, H., et al. (2011). The biomass mission: Mapping global forest biomass to better understand the terrestrial carbon cycle. *Remote Sensing of Environment*, 115(11):2850–2860.
- Næsset, E. (2011). Estimating above-ground biomass in young forests with airborne laser scanning. *International Journal of Remote Sensing*, 32(2):473–501.

- Neigh, C., Nelson, R., Ranson, K., Margolis, H., Montesano, P., Sun, G., Kharuk, V., Næsset, E., Wulder, M., and Andersen, H. (2013). Taking stock of circumboreal forest carbon with ground measurements, airborne and spaceborne lidar. *Remote Sensing of Environment*, 137:274–287.
- Ometto, J. P., Aguiar, A. P., Assis, T., Soler, L., Valle, P., Tejada, G., Lapola, D. M., and Meir, P. (2014). Amazon forest biomass density maps: Tackling the uncertainty in carbon emission estimates. *Climatic Change*, pages 1–16.
- Schimel, D. and Keller, M. (2015). Big questions, big science: meeting the challenges of global ecology. *Oecologia*, 177(4):925–934.
- UN-REDD (2009). The un-redd programme. <http://www.un-redd.org/>. Accessed: 5-22-2014.
- UNFCCC (2015). United nations framework convention on climate change. <http://unfccc.int/2860.php>. Accessed: 6-6-15.
- West, P. (2004). *Tree and Forest Measurement*. Springer.
- Wood, F., Kochenderfer, J. N., and Adams, M. B. (2016). Felled tree biomass for four hardwood species of the central appalachians, west virginia. fort collins, co: Forest service research data archive.